

Saniya Mihani

AI Engineer (Production Systems) · Agentic AI · LLM Infrastructure · GenAI MLOps
Mumbai, India · saniyamihani031@gmail.com · +91 7620953494 · LinkedIn · GitHub

Education

Thadomal Shahani Engineering College, Mumbai University
Bachelor of Engineering in Information Technology

Aug 2023 – Mar 2027

Experience

Freelance AI/Full-Stack Engineer, Prism HRC

Aug 2026 – Present

- Engaged as a contract engineer by Prism HRC to design and build an AI-powered resume/career platform for prismhrc.com that parses uploaded resumes and LinkedIn profiles into structured sections.
 - Architecting an OpenAI-powered agent that rewrites resume bullets into results-driven lines, tailors a resume to a pasted job description in one click, and scores it against ATS keyword-match and formatting checks.
 - Owning end-to-end delivery as the sole engineer on the project, from system architecture through deployment.
- Stack: Python, FastAPI, Next.js, LangGraph, OpenAI API, PostgreSQL*

Technical Skills

Languages: Python, TypeScript, JavaScript, SQL

AI / LLMs: Generative AI, RAG, LLM Fine-Tuning (LoRA, QLoRA), SFT, DPO, LLM Alignment, Agentic AI, Conversational Voice AI, Tool Calling, Structured Outputs, Prompt Engineering, Hallucination Mitigation, RAGAS Evaluation

Frameworks: LangChain, LangGraph, FastAPI, Transformers, Hugging Face, sentence-transformers, BGE Embeddings, Vapi, Groq, OpenAI API, Azure OpenAI, Weights & Biases

Cloud: AWS (ECS, ECR, ALB, API Gateway, CloudWatch), Azure (Container Apps, Container Registry, Application Insights), Vercel

Databases: PostgreSQL, Redis, Pinecone, ChromaDB

MLOps / DevOps: Docker, GitHub Actions CI/CD, Ruff, MyPy, Bandit, Pytest, OpenTelemetry, Langfuse, Structured Logging

Frontend / Infra: Next.js, React.js, WebSockets, SSE, Redis Pub/Sub, REST APIs, JWT, Git

Projects

FinGuard-APA: On-Premise Multi-Agent KYC/AML Compliance System

Present

- Architected a multi-agent compliance system with LangGraph to screen entities against OFAC, SEBI, and FinCEN sanctions lists, using hybrid BM25 + dense retrieval with cross-encoder reranking for fuzzy name matching and an explainable, citation-grounded decision trail.
- Built a self-critique guardrail and an LLM-as-judge evaluation harness on Mistral-7B (via Ollama), instrumented with Langfuse/OpenTelemetry observability and Selenium RPA integration, surfaced through a compliance-officer sign-off UI for human-in-the-loop review.
- Continuing domain-adaptive pretraining of a small open-weight model on compliance and financial-regulatory text, then adapting it for fuzzy sanctions-name matching, deployed via Ollama on the existing Mistral-7B serving layer.

Stack: LangGraph, ChromaDB, Mistral-7B (Ollama), Selenium RPA, Langfuse, OpenTelemetry

Loan Operation Intelligence: AI Voice Ops & Compliance Platform

Live Demo | GitHub

- Built a production outbound AI voice-bot platform for loan servicing, integrating grounded RAG retrieval with a real-time speech pipeline; achieved 98.7% retrieval accuracy and sub-850ms P95 ASR latency across live borrower calls.
- Designed a PII ingestion shield to scrub sensitive borrower data (phone, tax, email) pre-indexing, built multilingual call handling for Taglish and Bahasa Indonesia, and streamed live compliance/objection signals to a supervisor dashboard via Redis Pub/Sub and SSE.

Stack: Python, FastAPI, Next.js, Vapi, BGE Embeddings, Pinecone, PostgreSQL, Redis, Langfuse

ResearchOS: Production Agentic AI Research Platform

Live Demo | GitHub

- Architected a production multi-agent research platform on AWS (ECS + ALB + ECR) using LangGraph orchestration to coordinate modular retrieval, reasoning, citation-verification, and report-generation agents through event-driven SSE and WebSocket streaming pipelines.
- Engineered a FastAPI backend with JWT auth, PostgreSQL, Redis caching, and Pinecone vector search; deployed a Next.js frontend on Vercel with real-time agent execution timelines and CloudWatch observability.

Stack: Python, LangGraph, FastAPI, PostgreSQL, Redis, Pinecone, Docker, AWS (ECS/ECR/ALB/CloudWatch), Next.js, Vercel

Production GenAI MLOps Platform

Live Demo | GitHub

- Engineered a multi-provider LLM Gateway unifying Groq, OpenAI, and Azure OpenAI under a single FastAPI backend with intelligent routing, provider health monitoring, and OpenTelemetry distributed tracing.
- Automated end-to-end CI/CD with GitHub Actions (linting, type checks, security scans, tests, AI evaluation gates) and deployed to Azure Container Apps with staging/production separation; shipped a Next.js frontend with AI Playground and Health/Metrics/Evaluation dashboards.

Stack: Python, FastAPI, Docker, GitHub Actions, Azure (Container Apps, ACR, App Insights), OpenTelemetry, Groq, OpenAI, Next.js, Vercel

Certifications

RAG Systems, DeepLearning.AI

Jun 2026

Agentic AI, DeepLearning.AI

May 2026

Build Interactive Agents with Generative UI, DeepLearning.AI

May 2026